

The Epistemology of Computational Systems

A. R. Fontaine

M. J. Hessler

2024

Abstract

This paper examines the epistemological foundations of modern computational systems, focusing on the boundary conditions under which formal verification methods yield reliable knowledge claims. We argue that Gettier-style counterexamples apply with equal force to machine-generated proof certificates.¹ A taxonomy of computational epistemic failures is proposed, grounded in an analysis of three representative system classes.

1 Introduction

Computational epistemology sits at the intersection of formal methods, philosophy of science, and cognitive systems theory. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetur id, vulputate a,

magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum. A defining challenge² is distinguishing between *verified* results and *justified* results.

Nam dui ligula, fringilla a, euismod sodales, sollicitudin vel, wisi. Morbi auctor lorem non justo. Nam lacus libero, pretium at, lobortis vitae, ultricies et, tellus. Donec aliquet, tortor sed accumsan bibendum, erat ligula aliquet magna, vitae ornare odio metus a mi. Morbi ac orci et nisl hendrerit mollis. Suspendisse ut massa. Cras nec ante. Pellentesque

¹The term “proof certificate” is used here in the sense of Necula & Lee (1998), not the informal usage common in software engineering.

²This challenge was first articulated in the context of automated theorem provers by Robinson (1965), though he did not use this terminology.

tesque a nulla. Cum sociis natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Aliquam tincidunt urna. Nulla ullamcorper vestibulum turpis. Pellentesque cursus luctus mauris.

The distinction matters in practice: a system may produce a verified proof of a false theorem if the formal model diverges from the intended semantics.³ We return to this point in Section 2.

1.1 Scope and Method

The analysis is restricted to deterministic, finitely-computable systems. Probabilistic and quantum computational models are deferred to future work. Nulla malesuada porttitor diam. Donec felis erat, congue non, volutpat at, tincidunt tristique, libero. Vivamus viverra fermentum felis. Donec nuncummy pellentesque ante. Phasellus adipiscing semper elit. Proin fermentum massa ac quam. Sed diam turpis, molestie vitae, placerat a, molestie nec, leo. Maecenas lacinia. Nam ipsum ligula, eleifend at, accumsan nec, suscipit a, ipsum. Morbi blandit ligula feugiat magna. Nunc eleifend consequat lorem. Sed lacinia nulla vitae enim. Pellentesque tincidunt purus vel magna. Integer non enim. Praesent euismod nunc eu purus. Donec bibendum quam in tellus. Nullam cursus pulvinar lectus. Donec et mi. Nam vulputate metus eu enim. Vestibulum pellentesque felis eu massa.

³The canonical example is the Pentium FDIV bug (1994), in which a formally verified microcode routine produced incorrect results for a specific class of floating-point inputs.

1.1.1 Formal Preliminaries

Let Σ be a finite alphabet and $L \subseteq \Sigma^*$ a decidable language. A *verifier* V is a deterministic Turing machine such that for all $x \in L$, there exists a certificate c with $V(x, c) = 1$. Quisque ullamcorper placerat ipsum. Cras nibh. Morbi vel justo vitae lacus tincidunt ultrices. Lorem ipsum dolor sit amet, consectetur adipiscing elit. In hac habitasse platea dictumst. Integer tempus convallis augue. Etiam facilisis. Nunc elementum fermentum wisi. Aenean placerat. Ut imperdiet, enim sed gravida sollicitudin, felis odio placerat quam, ac pulvinar elit purus eget enim. Nunc vitae tortor. Proin tempus nibh sit amet nisl. Vivamus quis tortor vitae risus porta vehicula.

2 A Taxonomy of Epistemic Failure Modes

We identify four primary failure modes in computational knowledge systems. Each mode corresponds to a distinct breakdown in the justification chain from axioms to conclusions.

2.1 Type I: Model-Semantic Divergence

Model-semantic divergence occurs when the formal model M is a sound abstraction of the specification S but S is itself an inaccurate description of the intended system $\mathcal{I}$. Fusce mauris. Vestibulum luctus nibh at lectus. Sed bibendum, nulla a faucibus semper,

leo velit ultricies tellus, ac venenatis arcu wisi vel nisl. Vestibulum diam. Aliquam pellen-tesque, augue quis sagittis posuere, turpis la-cus congue quam, in hendrerit risus eros eget felis. Maecenas eget erat in sapien mattis porttitor. Vestibulum porttitor. Nulla faci-lisi. Sed a turpis eu lacus commodo facilisi. Morbi fringilla, wisi in dignissim interdum, justo lectus sagittis dui, et vehicula libero dui cursus dui. Mauris tempor ligula sed lacus. Duis cursus enim ut augue. Cras ac magna. Cras nulla. Nulla egestas. Curabitur a leo. Quisque egestas wisi eget nunc. Nam feugiat lacus vel est. Curabitur consectetur.

2.1.1 Illustrative Cases

The TLS 1.3 handshake verification effort [1] provides a well-documented instance of this failure mode.⁴ Suspendisse vel felis. Ut lo-rem lorem, interdum eu, tincidunt sit amet, laoreet vitae, arcu. Aenean faucibus pede eu ante. Praesent enim elit, rutrum at, mole-stie non, nonummy vel, nisl. Ut lectus eros, malesuada sit amet, fermentum eu, sodales cursus, magna. Donec eu purus. Quisque ve-hicula, urna sed ultricies auctor, pede lorem egestas dui, et convallis elit erat sed nulla. Donec luctus. Curabitur et nunc. Aliquam dolor odio, commodo pretium, ultricies non, pharetra in, velit. Integer arcu est, nonummy in, fermentum faucibus, egestas vel, odio.

⁴Bhargavan et al. verified the cryptographic core of TLS 1.3 using F $\star$, but noted that their model ex-cluded session resumption tickets—an exclusion that proved non-trivial in later implementations.

2.2 Type II: Proof-State Cor-ruption

Proof-state corruption arises when a proof assistant’s internal state is modified by a trusted kernel extension in a way that is not reflected in the proof certificate. Sed com-modo posuere pede. Mauris ut est. Ut quis purus. Sed ac odio. Sed vehicula hendre-rit sem. Duis non odio. Morbi ut dui. Sed accumsan risus eget odio. In hac habitasse platea dictumst. Pellentesque non elit. Fu-sce sed justo eu urna porta tincidunt. Mauris felis odio, sollicitudin sed, volutpat a, ornare ac, erat. Morbi quis dolor. Donec pellen-tesque, erat ac sagittis semper, nunc dui lobortis purus, quis congue purus metus ultricies tel-lus. Proin et quam. Class aptent taciti socio-squ ad litora torquent per conubia nostra, per inceptos hymenaeos. Praesent sapien turpis, fermentum vel, eleifend faucibus, vehicula eu, lacus.




Figure 1: Schematic of the justification chain from axioms ($\mathcal{A}$) through the proof kernel (K) to the certificate (C). A Type II failure corrupts the $K \rightarrow C$ link without invalidat-ing $\mathcal{A} \rightarrow K$.

Pellentesque habitant morbi tristique se-nectus et netus et malesuada fames ac turpis

egestas. Donec odio elit, dictum in, hendrerit sit amet, egestas sed, leo. Praesent feugiat sapien aliquet odio. Integer vitae justo. Aliquam vestibulum fringilla lorem. Sed neque lectus, consectetur at, consectetur sed, eleifend ac, lectus. Nulla facilisi. Pellentesque eget lectus. Proin eu metus. Sed porttitor. In hac habitasse platea dictumst. Suspendisse eu lectus. Ut mi mi, lacinia sit amet, placerat et, mollis vitae, dui. Sed ante tellus, tristique ut, iaculis eu, malesuada ac, dui. Mauris nibh leo, facilisis non, adipiscing quis, ultrices a, dui.

2.3 Type III: Certificate Non-Portability

A certificate c produced by verifier V_1 may not be accepted by verifier V_2 even when both are sound with respect to the same logic, because the two systems may use non-equivalent encodings of the proof term. Morbi luctus, wisi viverra faucibus pretium, nibh est placerat odio, nec commodo wisi enim eget quam. Quisque libero justo, consectetur a, feugiat vitae, porttitor eu, libero. Suspendisse sed mauris vitae elit sollicitudin malesuada. Maecenas ultricies eros sit amet ante. Ut venenatis velit. Maecenas sed mi eget dui varius euismod. Phasellus aliquet volutpat odio. Vestibulum ante ipsum primis in faucibus orci luctus et ultrices posuere cubilia Curae; Pellentesque sit amet pede ac sem eleifend consectetur. Nullam elementum, urna vel imperdiet sodales, elit ipsum pharetra ligula, ac pretium ante justo a nulla. Curabitur tristique arcu eu metus. Vestibu-

lum lectus. Proin mauris. Proin eu nunc eu urna hendrerit faucibus. Aliquam auctor, pede consequat laoreet varius, eros tellus scelerisque quam, pellentesque hendrerit ipsum dolor sed augue. Nulla nec lacus.

2.3.1 Cross-System Verification

Cross-system verification remains an open problem.⁵ Suspendisse vitae elit. Aliquam arcu neque, ornare in, ullamcorper quis, commodo eu, libero. Fusce sagittis erat at erat tristique mollis. Maecenas sapien libero, molestie et, lobortis in, sodales eget, dui. Morbi ultrices rutrum lorem. Nam elementum ullamcorper leo. Morbi dui. Aliquam sagittis. Nunc placerat. Pellentesque tristique sodales est. Maecenas imperdiet lacinia velit. Cras non urna. Morbi eros pede, suscipit ac, varius vel, egestas non, eros. Praesent malesuada, diam id pretium elementum, eros sem dictum tortor, vel consectetur odio sem sed wisi.

3 Implications for Verified Software Engineering

The taxonomy developed in Section 2 has direct implications for the deployment of formally verified components in safety-critical systems.

⁵The QED Manifesto (1994) proposed a universal proof repository, but the project stalled precisely because of certificate non-portability across the Coq, Isabelle, and HOL family of proof assistants.

3.1 Certification Standards

Current certification standards such as DO-178C (avionics) and IEC 61508 (industrial safety) do not yet account for Type II or Type III failures. Sed feugiat. Cum sociis natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Ut pellentesque augue sed urna. Vestibulum diam eros, fringilla et, consectetur eu, nonummy id, sapien. Nullam at lectus. In sagittis ultrices mauris. Curabitur malesuada erat sit amet massa. Fusce blandit. Aliquam erat volutpat. Aliquam euismod. Aenean vel lectus. Nunc imperdiet justo nec dolor.

3.1.1 Proposed Augmentations

We propose three augmentations to existing standards. First, certificate portability testing should be mandated for all Level A software.⁶ Second, formal models should include explicit semantic gap documentation. Third, trusted kernel extensions should undergo independent formal verification.

Etiam euismod. Fusce facilisis lacinia dui. Suspendisse potenti. In mi erat, cursus id, nonummy sed, ullamcorper eget, sapien. Praesent pretium, magna in eleifend egestas, pede pede pretium lorem, quis consectetur tortor sapien facilisis magna. Mauris quis magna varius nulla scelerisque imperdiet. Aliquam non quam. Aliquam porttitor quam a lacus. Praesent vel arcu ut tortor cursus volutpat. In vitae pede quis diam bi-

⁶Level A software is defined by DO-178C as software whose failure would result in catastrophic failure conditions.

bendum placerat. Fusce elementum convallis neque. Sed dolor orci, scelerisque ac, dapibus nec, ultricies ut, mi. Duis nec dui quis leo sagittis commodo.

3.2 Tool Qualification

Tool qualification under DO-330 partially addresses Type I failures but does not cover the certificate layer. Aliquam lectus. Vivamus leo. Quisque ornare tellus ullamcorper nulla. Mauris porttitor pharetra tortor. Sed fringilla justo sed mauris. Mauris tellus. Sed non leo. Nullam elementum, magna in cursus sodales, augue est scelerisque sapien, venenatis congue nulla arcu et pede. Ut suscipit enim vel sapien. Donec congue. Maecenas urna mi, suscipit in, placerat ut, vestibulum ut, massa. Fusce ultrices nulla et nisl.

4 Conclusion

We have presented a taxonomy of four epistemic failure modes in computational verification systems and argued that existing certification standards are insufficient to address the full range of failure modes. Etiam ac leo a risus tristique nonummy. Donec dignissim tincidunt nulla. Vestibulum rhoncus molestie odio. Sed lobortis, justo et pretium lobortis, mauris turpis condimentum augue, nec ultricies nibh arcu pretium enim. Nunc purus neque, placerat id, imperdiet sed, pellentesque nec, nisl. Vestibulum imperdiet neque non sem accumsan laoreet. In hac habitasse platea dictumst. Etiam condimentum facilisis libero. Suspendisse in elit quis nisl aliquam da-

pibus. Pellentesque auctor sapien. Sed egestas sapien nec lectus. Pellentesque vel dui vel neque bibendum viverra. Aliquam portitor nisl nec pede. Proin mattis libero vel turpis. Donec rutrum mauris et libero. Proin euismod porta felis. Nam lobortis, metus quis elementum commodo, nunc lectus elementum mauris, eget vulputate ligula tellus eu neque. Vivamus eu dolor.

References

- [1] K. Bhargavan, B. Blanchet, and N. Kobeissi. Verified models and reference implementations for the TLS 1.3 standard candidate. In *Proceedings of the IEEE Symposium on Security and Privacy*, pages 483–502, 2017.
- [2] G. C. Necula and P. Lee. Safe kernel extensions without run-time checking. In *Proceedings of the 2nd USENIX Symposium on Operating Systems Design and Implementation*, 1998.
- [3] J. A. Robinson. A machine-oriented logic based on the resolution principle. *Journal of the ACM*, 12(1):23–41, 1965.